

Transfer_learning_with_MobileNet_v1

October 31, 2025

1 Transfer Learning with MobileNetV2

Welcome to this week's assignment, where you'll be using transfer learning on a pre-trained CNN to build an Alpaca/Not Alpaca classifier!

A pre-trained model is a network that's already been trained on a large dataset and saved, which allows you to use it to customize your own model cheaply and efficiently. The one you'll be using, MobileNetV2, was designed to provide fast and computationally efficient performance. It's been pre-trained on ImageNet, a dataset containing over 14 million images and 1000 classes.

By the end of this assignment, you will be able to:

- Create a dataset from a directory
- Preprocess and augment data using the Sequential API
- Adapt a pretrained model to new data and train a classifier using the Functional API and MobileNet
- Fine-tune a classifier's final layers to improve accuracy

1.1 Important Note on Submission to the AutoGrader

Before submitting your assignment to the AutoGrader, please make sure you are not doing the following:

1. You have not added any *extra* **print** statement(s) in the assignment.
2. You have not added any *extra* code cell(s) in the assignment.
3. You have not changed any of the function parameters.
4. You are not using any global variables inside your graded exercises. Unless specifically instructed to do so, please refrain from it and use the local variables instead.
5. You are not changing the assignment code where it is not required, like creating *extra* variables.

If you do any of the following, you will get something like, **Grader Error: Grader feedback not found** (or similarly unexpected) error upon submitting your assignment. Before asking for help/debugging the errors in your assignment, check for these first. If this is the case, and you don't remember the changes you have made, you can get a fresh copy of the assignment by following these [instructions](#).

1.2 Table of Content

- Section ??

- Section ??
- Section ??
 - Section ??
- Section ??
 - Section ??
 - Section ??
 - * Section ??
 - Section ??
 - * Section ??

1 - Packages

```
[1]: ### v2.1
```

```
[2]: import matplotlib.pyplot as plt
import json
import numpy as np
import os
import tensorflow as tf
import tensorflow.keras.layers as tfl

from tensorflow.keras.preprocessing import image_dataset_from_directory
from tensorflow.keras.layers.experimental.preprocessing import RandomFlip, RandomRotation
```

1.1 Create the Dataset and Split it into Training and Validation Sets

When training and evaluating deep learning models in Keras, generating a dataset from image files stored on disk is simple and fast. Call `image_data_set_from_directory()` to read from the directory and create both training and validation datasets.

If you're specifying a validation split, you'll also need to specify the subset for each portion. Just set the training set to `subset='training'` and the validation set to `subset='validation'`.

You'll also set your seeds to match each other, so your training and validation sets don't overlap. :)

```
[3]: BATCH_SIZE = 32
IMG_SIZE = (160, 160)
directory = "dataset/"
train_dataset = image_dataset_from_directory(directory,
                                             shuffle=True,
                                             batch_size=BATCH_SIZE,
                                             image_size=IMG_SIZE,
                                             validation_split=0.2,
                                             subset='training',
                                             seed=42)

validation_dataset = image_dataset_from_directory(directory,
                                                  shuffle=True,
                                                  batch_size=BATCH_SIZE,
```

```
image_size=IMG_SIZE,  
validation_split=0.2,  
subset='validation',  
seed=42)
```

Found 327 files belonging to 2 classes.
Using 262 files for training.
Found 327 files belonging to 2 classes.
Using 65 files for validation.

Now let's take a look at some of the images from the training set:

Note: The original dataset has some mislabelled images in it as well.

```
[4]: class_names = train_dataset.class_names  
  
plt.figure(figsize=(10, 10))  
for images, labels in train_dataset.take(1):  
    for i in range(9):  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(images[i].numpy().astype("uint8"))  
        plt.title(class_names[labels[i]])  
        plt.axis("off")
```



2 - Preprocess and Augment Training Data

You may have encountered `dataset.prefetch` in a previous TensorFlow assignment, as an important extra step in data preprocessing.

Using `prefetch()` prevents a memory bottleneck that can occur when reading from disk. It sets aside some data and keeps it ready for when it's needed, by creating a source dataset from your input data, applying a transformation to preprocess it, then iterating over the dataset one element at a time. Because the iteration is streaming, the data doesn't need to fit into memory.

You can set the number of elements to prefetch manually, or you can use `tf.data.experimental.AUTOTUNE` to choose the parameters automatically. Autotune prompts `tf.data` to tune that value dynamically at runtime, by tracking the time spent in each operation and feeding those times into an optimization algorithm. The optimization algorithm tries to find

the best allocation of its CPU budget across all tunable operations.

To increase diversity in the training set and help your model learn the data better, it's standard practice to augment the images by transforming them, i.e., randomly flipping and rotating them. Keras' Sequential API offers a straightforward method for these kinds of data augmentations, with built-in, customizable preprocessing layers. These layers are saved with the rest of your model and can be re-used later. Ahh, so convenient!

As always, you're invited to read the official docs, which you can find for data augmentation [here](#).

```
[5]: AUTOTUNE = tf.data.experimental.AUTOTUNE
train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)
```

Exercise 1 - data_augmenter

Implement a function for data augmentation. Use a `Sequential` keras model composed of 2 layers:
* `RandomFlip('horizontal')` * `RandomRotation(0.2)`

```
[6]: # UNQ_C1
# GRADED FUNCTION: data_augmenter
def data_augmenter():
    """
    Create a Sequential model composed of 2 layers
    Returns:
        tf.keras.Sequential
    """
    ### START CODE HERE

    data_augmentation = tf.keras.Sequential()
    data_augmentation.add(RandomFlip('horizontal'))
    data_augmentation.add(RandomRotation(0.2))

    ### END CODE HERE

    return data_augmentation
```

```
[7]: augmenter = data_augmenter()

assert(augmenter.layers[0].name.startswith('random_flip')), "First layer must_
↳be RandomFlip"
assert augmenter.layers[0].mode == 'horizontal', "RadomFlip parameter must be_
↳horizontal"
assert(augmenter.layers[1].name.startswith('random_rotation')), "Second layer_
↳must be RandomRotation"
assert augmenter.layers[1].factor == 0.2, "Rotation factor must be 0.2"
assert len(augmenter.layers) == 2, "The model must have only 2 layers"

print('\033[92mAll tests passed!')
```



All tests passed!

Take a look at how an image from the training set has been augmented with simple transformations:

From one cute animal, to 9 variations of that cute animal, in three lines of code. Now your model has a lot more to learn from.

```
[8]: data_augmentation = data_augmenter()

for image, _ in train_dataset.take(1):
    plt.figure(figsize=(10, 10))
    first_image = image[0]
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        augmented_image = data_augmentation(tf.expand_dims(first_image, 0))
        plt.imshow(augmented_image[0] / 255)
        plt.axis('off')
```



Next, you'll apply your first tool from the MobileNet application in TensorFlow, to normalize your input. Since you're using a pre-trained model that was trained on the normalization values $[-1,1]$, it's best practice to reuse that standard with `tf.keras.applications.mobilenet_v2.preprocess_input`.

What you should remember:

- When calling `image_data_set_from_directory()`, specify the train/val subsets and match the seeds to prevent overlap
- Use `prefetch()` to prevent memory bottlenecks when reading from disk
- Give your model more to learn from with simple data augmentations like rotation and flipping.
- When using a pretrained model, it's best to reuse the weights it was trained on.

```
[9]: preprocess_input = tf.keras.applications.mobilenet_v2.preprocess_input
```

3 - Using MobileNetV2 for Transfer Learning

MobileNetV2 was trained on ImageNet and is optimized to run on mobile and other low-power applications. It's 155 layers deep (just in case you felt the urge to plot the model yourself, prepare for a long journey!) and very efficient for object detection and image segmentation tasks, as well as classification tasks like this one. The architecture has three defining characteristics:

- Depthwise separable convolutions
- Thin input and output bottlenecks between layers
- Shortcut connections between bottleneck layers

3.1 - Inside a MobileNetV2 Convolutional Building Block

MobileNetV2 uses depthwise separable convolutions as efficient building blocks. Traditional convolutions are often very resource-intensive, and depthwise separable convolutions are able to reduce the number of trainable parameters and operations and also speed up convolutions in two steps:

1. The first step calculates an intermediate result by convolving on each of the channels independently. This is the depthwise convolution.
2. In the second step, another convolution merges the outputs of the previous step into one. This gets a single result from a single feature at a time, and then is applied to all the filters in the output layer. This is the pointwise convolution, or: **Shape of the depthwise convolution X Number of filters.**

Figure 1 : MobileNetV2 Architecture This diagram was inspired by the original seen here.

Each block consists of an inverted residual structure with a bottleneck at each end. These bottlenecks encode the intermediate inputs and outputs in a low dimensional space, and prevent non-linearities from destroying important information.

The shortcut connections, which are similar to the ones in traditional residual networks, serve the same purpose of speeding up training and improving predictions. These connections skip over the intermediate convolutions and connect the bottleneck layers.

Let's try to train your base model using all the layers from the pretrained model.

Similarly to how you reused the pretrained normalization values MobileNetV2 was trained on, you'll also load the pretrained weights from ImageNet.

By specifying `weights='imagenet'`, the weights are downloaded for ImageNet. But for this lab, you'll load them locally from your workspace.

```
[10]: IMG_SHAPE = IMG_SIZE + (3,)
base_model_path="imagenet_base_model/
↳with_top_mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_160.h5"

base_model = tf.keras.applications.MobileNetV2(input_shape=IMG_SHAPE,
                                                include_top=True,
                                                weights=base_model_path)
```

Print the model summary below to see all the model's layers, the shapes of their outputs, and the total number of parameters, trainable and non-trainable.

```
[11]: base_model.summary()
```

Model: "mobilenetv2_1.00_160"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 160, 160, 3)]	0	
Conv1_pad (ZeroPadding2D)	(None, 161, 161, 3)	0	input_1[0][0]
Conv1 (Conv2D)	(None, 80, 80, 32)	864	Conv1_pad[0][0]
bn_Conv1 (BatchNormalization)	(None, 80, 80, 32)	128	Conv1[0][0]
Conv1_relu (ReLU)	(None, 80, 80, 32)	0	bn_Conv1[0][0]
expanded_conv_depthwise (Depthwise Conv2D)	(None, 80, 80, 32)	288	Conv1_relu[0][0]
expanded_conv_depthwise_BN (BatchNormalization)	(None, 80, 80, 32)	128	expanded_conv_depthwise[0][0]
expanded_conv_depthwise_relu (ReLU)	(None, 80, 80, 32)	0	


```

expanded_conv_depthwise_BN[0][0]
-----
expanded_conv_project (Conv2D) (None, 80, 80, 16) 512
expanded_conv_depthwise_relu[0][0]
-----
expanded_conv_project_BN (Batch (None, 80, 80, 16) 64
expanded_conv_project[0][0]
-----
block_1_expand (Conv2D) (None, 80, 80, 96) 1536
expanded_conv_project_BN[0][0]
-----
block_1_expand_BN (BatchNormali (None, 80, 80, 96) 384
block_1_expand[0][0]
-----
block_1_expand_relu (ReLU) (None, 80, 80, 96) 0
block_1_expand_BN[0][0]
-----
block_1_pad (ZeroPadding2D) (None, 81, 81, 96) 0
block_1_expand_relu[0][0]
-----
block_1_depthwise (DepthwiseCon (None, 40, 40, 96) 864
block_1_pad[0][0]
-----
block_1_depthwise_BN (BatchNorm (None, 40, 40, 96) 384
block_1_depthwise[0][0]
-----
block_1_depthwise_relu (ReLU) (None, 40, 40, 96) 0
block_1_depthwise_BN[0][0]
-----
block_1_project (Conv2D) (None, 40, 40, 24) 2304
block_1_depthwise_relu[0][0]
-----
block_1_project_BN (BatchNormal (None, 40, 40, 24) 96
block_1_project[0][0]
-----
block_2_expand (Conv2D) (None, 40, 40, 144) 3456

```

block_1_project_BN[0][0]

block_2_expand_BN (BatchNormali (None, 40, 40, 144) 576
block_2_expand[0][0]

block_2_expand_relu (ReLU) (None, 40, 40, 144) 0
block_2_expand_BN[0][0]

block_2_depthwise (DepthwiseCon (None, 40, 40, 144) 1296
block_2_expand_relu[0][0]

block_2_depthwise_BN (BatchNorm (None, 40, 40, 144) 576
block_2_depthwise[0][0]

block_2_depthwise_relu (ReLU) (None, 40, 40, 144) 0
block_2_depthwise_BN[0][0]

block_2_project (Conv2D) (None, 40, 40, 24) 3456
block_2_depthwise_relu[0][0]

block_2_project_BN (BatchNormal (None, 40, 40, 24) 96
block_2_project[0][0]

block_2_add (Add) (None, 40, 40, 24) 0
block_1_project_BN[0][0]
block_2_project_BN[0][0]

block_3_expand (Conv2D) (None, 40, 40, 144) 3456
block_2_add[0][0]

block_3_expand_BN (BatchNormali (None, 40, 40, 144) 576
block_3_expand[0][0]

block_3_expand_relu (ReLU) (None, 40, 40, 144) 0
block_3_expand_BN[0][0]

```

block_3_pad (ZeroPadding2D)      (None, 41, 41, 144)  0
block_3_expand_relu[0][0]
-----

block_3_depthwise (DepthwiseCon (None, 20, 20, 144)  1296
block_3_pad[0][0]
-----

block_3_depthwise_BN (BatchNorm (None, 20, 20, 144)  576
block_3_depthwise[0][0]
-----

block_3_depthwise_relu (ReLU)    (None, 20, 20, 144)  0
block_3_depthwise_BN[0][0]
-----

block_3_project (Conv2D)         (None, 20, 20, 32)   4608
block_3_depthwise_relu[0][0]
-----

block_3_project_BN (BatchNormal (None, 20, 20, 32)   128
block_3_project[0][0]
-----

block_4_expand (Conv2D)         (None, 20, 20, 192)  6144
block_3_project_BN[0][0]
-----

block_4_expand_BN (BatchNormali (None, 20, 20, 192)  768
block_4_expand[0][0]
-----

block_4_expand_relu (ReLU)      (None, 20, 20, 192)  0
block_4_expand_BN[0][0]
-----

block_4_depthwise (DepthwiseCon (None, 20, 20, 192)  1728
block_4_expand_relu[0][0]
-----

block_4_depthwise_BN (BatchNorm (None, 20, 20, 192)  768
block_4_depthwise[0][0]
-----

block_4_depthwise_relu (ReLU)   (None, 20, 20, 192)  0
block_4_depthwise_BN[0][0]
-----

```

block_4_project (Conv2D) (None, 20, 20, 32) 6144
block_4_depthwise_relu[0][0]

block_4_project_BN (BatchNormal (None, 20, 20, 32) 128
block_4_project[0][0]

block_4_add (Add) (None, 20, 20, 32) 0
block_3_project_BN[0][0]
block_4_project_BN[0][0]

block_5_expand (Conv2D) (None, 20, 20, 192) 6144
block_4_add[0][0]

block_5_expand_BN (BatchNormali (None, 20, 20, 192) 768
block_5_expand[0][0]

block_5_expand_relu (ReLU) (None, 20, 20, 192) 0
block_5_expand_BN[0][0]

block_5_depthwise (DepthwiseCon (None, 20, 20, 192) 1728
block_5_expand_relu[0][0]

block_5_depthwise_BN (BatchNorm (None, 20, 20, 192) 768
block_5_depthwise[0][0]

block_5_depthwise_relu (ReLU) (None, 20, 20, 192) 0
block_5_depthwise_BN[0][0]

block_5_project (Conv2D) (None, 20, 20, 32) 6144
block_5_depthwise_relu[0][0]

block_5_project_BN (BatchNormal (None, 20, 20, 32) 128
block_5_project[0][0]

block_5_add (Add) (None, 20, 20, 32) 0
block_4_add[0][0]
block_5_project_BN[0][0]

```

-----
-----
block_6_expand (Conv2D)          (None, 20, 20, 192) 6144
block_5_add[0][0]
-----
-----
block_6_expand_BN (BatchNormali (None, 20, 20, 192) 768
block_6_expand[0][0]
-----
-----
block_6_expand_relu (ReLU)       (None, 20, 20, 192) 0
block_6_expand_BN[0][0]
-----
-----
block_6_pad (ZeroPadding2D)      (None, 21, 21, 192) 0
block_6_expand_relu[0][0]
-----
-----
block_6_depthwise (DepthwiseCon (None, 10, 10, 192) 1728
block_6_pad[0][0]
-----
-----
block_6_depthwise_BN (BatchNorm (None, 10, 10, 192) 768
block_6_depthwise[0][0]
-----
-----
block_6_depthwise_relu (ReLU)    (None, 10, 10, 192) 0
block_6_depthwise_BN[0][0]
-----
-----
block_6_project (Conv2D)         (None, 10, 10, 64) 12288
block_6_depthwise_relu[0][0]
-----
-----
block_6_project_BN (BatchNormal (None, 10, 10, 64) 256
block_6_project[0][0]
-----
-----
block_7_expand (Conv2D)          (None, 10, 10, 384) 24576
block_6_project_BN[0][0]
-----
-----
block_7_expand_BN (BatchNormali (None, 10, 10, 384) 1536
block_7_expand[0][0]
-----
-----
block_7_expand_relu (ReLU)       (None, 10, 10, 384) 0
block_7_expand_BN[0][0]

```

```

-----
-----
block_7_depthwise (DepthwiseCon (None, 10, 10, 384) 3456
block_7_expand_relu[0][0]
-----
-----
block_7_depthwise_BN (BatchNorm (None, 10, 10, 384) 1536
block_7_depthwise[0][0]
-----
-----
block_7_depthwise_relu (ReLU) (None, 10, 10, 384) 0
block_7_depthwise_BN[0][0]
-----
-----
block_7_project (Conv2D) (None, 10, 10, 64) 24576
block_7_depthwise_relu[0][0]
-----
-----
block_7_project_BN (BatchNormal (None, 10, 10, 64) 256
block_7_project[0][0]
-----
-----
block_7_add (Add) (None, 10, 10, 64) 0
block_6_project_BN[0][0]
block_7_project_BN[0][0]
-----
-----
block_8_expand (Conv2D) (None, 10, 10, 384) 24576
block_7_add[0][0]
-----
-----
block_8_expand_BN (BatchNormali (None, 10, 10, 384) 1536
block_8_expand[0][0]
-----
-----
block_8_expand_relu (ReLU) (None, 10, 10, 384) 0
block_8_expand_BN[0][0]
-----
-----
block_8_depthwise (DepthwiseCon (None, 10, 10, 384) 3456
block_8_expand_relu[0][0]
-----
-----
block_8_depthwise_BN (BatchNorm (None, 10, 10, 384) 1536
block_8_depthwise[0][0]
-----
-----
block_8_depthwise_relu (ReLU) (None, 10, 10, 384) 0

```


block_8_depthwise_BN[0][0]

block_8_project (Conv2D) (None, 10, 10, 64) 24576
block_8_depthwise_relu[0][0]

block_8_project_BN (BatchNormal (None, 10, 10, 64) 256
block_8_project[0][0]

block_8_add (Add) (None, 10, 10, 64) 0
block_7_add[0][0]
block_8_project_BN[0][0]

block_9_expand (Conv2D) (None, 10, 10, 384) 24576
block_8_add[0][0]

block_9_expand_BN (BatchNormali (None, 10, 10, 384) 1536
block_9_expand[0][0]

block_9_expand_relu (ReLU) (None, 10, 10, 384) 0
block_9_expand_BN[0][0]

block_9_depthwise (DepthwiseCon (None, 10, 10, 384) 3456
block_9_expand_relu[0][0]

block_9_depthwise_BN (BatchNorm (None, 10, 10, 384) 1536
block_9_depthwise[0][0]

block_9_depthwise_relu (ReLU) (None, 10, 10, 384) 0
block_9_depthwise_BN[0][0]

block_9_project (Conv2D) (None, 10, 10, 64) 24576
block_9_depthwise_relu[0][0]

block_9_project_BN (BatchNormal (None, 10, 10, 64) 256
block_9_project[0][0]

block_9_add (Add)	(None, 10, 10, 64)	0
block_8_add[0][0]		
block_9_project_BN[0][0]		

block_10_expand (Conv2D)	(None, 10, 10, 384)	24576
block_9_add[0][0]		

block_10_expand_BN (BatchNormal	(None, 10, 10, 384)	1536
block_10_expand[0][0]		

block_10_expand_relu (ReLU)	(None, 10, 10, 384)	0
block_10_expand_BN[0][0]		

block_10_depthwise (DepthwiseCo	(None, 10, 10, 384)	3456
block_10_expand_relu[0][0]		

block_10_depthwise_BN (BatchNor	(None, 10, 10, 384)	1536
block_10_depthwise[0][0]		

block_10_depthwise_relu (ReLU)	(None, 10, 10, 384)	0
block_10_depthwise_BN[0][0]		

block_10_project (Conv2D)	(None, 10, 10, 96)	36864
block_10_depthwise_relu[0][0]		

block_10_project_BN (BatchNorma	(None, 10, 10, 96)	384
block_10_project[0][0]		

block_11_expand (Conv2D)	(None, 10, 10, 576)	55296
block_10_project_BN[0][0]		

block_11_expand_BN (BatchNormal	(None, 10, 10, 576)	2304
block_11_expand[0][0]		

block_11_expand_relu (ReLU)	(None, 10, 10, 576)	0
block_11_expand_BN[0][0]		

block_11_depthwise (DepthwiseCo (None, 10, 10, 576) 5184
block_11_expand_relu[0][0]

block_11_depthwise_BN (BatchNor (None, 10, 10, 576) 2304
block_11_depthwise[0][0]

block_11_depthwise_relu (ReLU) (None, 10, 10, 576) 0
block_11_depthwise_BN[0][0]

block_11_project (Conv2D) (None, 10, 10, 96) 55296
block_11_depthwise_relu[0][0]

block_11_project_BN (BatchNorma (None, 10, 10, 96) 384
block_11_project[0][0]

block_11_add (Add) (None, 10, 10, 96) 0
block_10_project_BN[0][0]
block_11_project_BN[0][0]

block_12_expand (Conv2D) (None, 10, 10, 576) 55296
block_11_add[0][0]

block_12_expand_BN (BatchNormal (None, 10, 10, 576) 2304
block_12_expand[0][0]

block_12_expand_relu (ReLU) (None, 10, 10, 576) 0
block_12_expand_BN[0][0]

block_12_depthwise (DepthwiseCo (None, 10, 10, 576) 5184
block_12_expand_relu[0][0]

block_12_depthwise_BN (BatchNor (None, 10, 10, 576) 2304
block_12_depthwise[0][0]

block_12_depthwise_relu (ReLU) (None, 10, 10, 576) 0
block_12_depthwise_BN[0][0]

```

-----
-----
block_12_project (Conv2D)          (None, 10, 10, 96)    55296
block_12_depthwise_relu[0][0]
-----
-----
block_12_project_BN (BatchNorma (None, 10, 10, 96)    384
block_12_project[0][0]
-----
-----
block_12_add (Add)                  (None, 10, 10, 96)    0
block_11_add[0][0]
block_12_project_BN[0][0]
-----
-----
block_13_expand (Conv2D)           (None, 10, 10, 576)   55296
block_12_add[0][0]
-----
-----
block_13_expand_BN (BatchNormal (None, 10, 10, 576)   2304
block_13_expand[0][0]
-----
-----
block_13_expand_relu (ReLU)        (None, 10, 10, 576)   0
block_13_expand_BN[0][0]
-----
-----
block_13_pad (ZeroPadding2D)       (None, 11, 11, 576)   0
block_13_expand_relu[0][0]
-----
-----
block_13_depthwise (DepthwiseCo (None, 5, 5, 576)    5184
block_13_pad[0][0]
-----
-----
block_13_depthwise_BN (BatchNor (None, 5, 5, 576)    2304
block_13_depthwise[0][0]
-----
-----
block_13_depthwise_relu (ReLU)     (None, 5, 5, 576)     0
block_13_depthwise_BN[0][0]
-----
-----
block_13_project (Conv2D)          (None, 5, 5, 160)     92160
block_13_depthwise_relu[0][0]
-----
-----
block_13_project_BN (BatchNorma (None, 5, 5, 160)     640

```

block_13_project[0][0]

block_14_expand (Conv2D) (None, 5, 5, 960) 153600
block_13_project_BN[0][0]

block_14_expand_BN (BatchNormal (None, 5, 5, 960) 3840
block_14_expand[0][0]

block_14_expand_relu (ReLU) (None, 5, 5, 960) 0
block_14_expand_BN[0][0]

block_14_depthwise (DepthwiseCo (None, 5, 5, 960) 8640
block_14_expand_relu[0][0]

block_14_depthwise_BN (BatchNor (None, 5, 5, 960) 3840
block_14_depthwise[0][0]

block_14_depthwise_relu (ReLU) (None, 5, 5, 960) 0
block_14_depthwise_BN[0][0]

block_14_project (Conv2D) (None, 5, 5, 160) 153600
block_14_depthwise_relu[0][0]

block_14_project_BN (BatchNorma (None, 5, 5, 160) 640
block_14_project[0][0]

block_14_add (Add) (None, 5, 5, 160) 0
block_13_project_BN[0][0]
block_14_project_BN[0][0]

block_15_expand (Conv2D) (None, 5, 5, 960) 153600
block_14_add[0][0]

block_15_expand_BN (BatchNormal (None, 5, 5, 960) 3840
block_15_expand[0][0]

block_15_expand_relu (ReLU)	(None, 5, 5, 960)	0
block_15_expand_BN[0][0]		

block_15_depthwise (DepthwiseCo	(None, 5, 5, 960)	8640
block_15_expand_relu[0][0]		

block_15_depthwise_BN (BatchNor	(None, 5, 5, 960)	3840
block_15_depthwise[0][0]		

block_15_depthwise_relu (ReLU)	(None, 5, 5, 960)	0
block_15_depthwise_BN[0][0]		

block_15_project (Conv2D)	(None, 5, 5, 160)	153600
block_15_depthwise_relu[0][0]		

block_15_project_BN (BatchNorma	(None, 5, 5, 160)	640
block_15_project[0][0]		

block_15_add (Add)	(None, 5, 5, 160)	0
block_14_add[0][0]		
block_15_project_BN[0][0]		

block_16_expand (Conv2D)	(None, 5, 5, 960)	153600
block_15_add[0][0]		

block_16_expand_BN (BatchNormal	(None, 5, 5, 960)	3840
block_16_expand[0][0]		

block_16_expand_relu (ReLU)	(None, 5, 5, 960)	0
block_16_expand_BN[0][0]		

block_16_depthwise (DepthwiseCo	(None, 5, 5, 960)	8640
block_16_expand_relu[0][0]		

block_16_depthwise_BN (BatchNor	(None, 5, 5, 960)	3840
block_16_depthwise[0][0]		

```

-----
block_16_depthwise_relu (ReLU) (None, 5, 5, 960) 0
block_16_depthwise_BN[0][0]
-----
block_16_project (Conv2D) (None, 5, 5, 320) 307200
block_16_depthwise_relu[0][0]
-----
block_16_project_BN (BatchNorma (None, 5, 5, 320) 1280
block_16_project[0][0]
-----
Conv_1 (Conv2D) (None, 5, 5, 1280) 409600
block_16_project_BN[0][0]
-----
Conv_1_bn (BatchNormalization) (None, 5, 5, 1280) 5120 Conv_1[0][0]
-----
out_relu (ReLU) (None, 5, 5, 1280) 0 Conv_1_bn[0][0]
-----
global_average_pooling2d (Globa (None, 1280) 0 out_relu[0][0]
-----
predictions (Dense) (None, 1000) 1281000
global_average_pooling2d[0][0]
=====
Total params: 3,538,984
Trainable params: 3,504,872
Non-trainable params: 34,112
-----

```

Note the last 2 layers here. They are the so called top layers, and they are responsible of the classification in the model

```
[12]: nb_layers = len(base_model.layers)
print(base_model.layers[nb_layers - 2].name)
print(base_model.layers[nb_layers - 1].name)
```

```
global_average_pooling2d
predictions
```

Notice some of the layers in the summary like Conv2D and DepthwiseConv2D and how they follow the progression of expansion to depthwise convolution to projection. In combination with BatchNormalization and ReLU, these make up the bottleneck layers mentioned earlier.

What you should remember:

- MobileNetV2's unique features are:
 - Depthwise separable convolutions that provide lightweight feature filtering and creation
 - Input and output bottlenecks that preserve important information on either end of the block
- Depthwise separable convolutions deal with both spatial and depth (number of channels) dimensions

Next, choose the first batch from the tensorflow dataset to use the images, and run it through the MobileNetV2 base model to test out the predictions on some of your images.

```
[13]: image_batch, label_batch = next(iter(train_dataset))
      feature_batch = base_model(image_batch)
      print(feature_batch.shape)
```

(32, 1000)

```
[14]: #Shows the different label probabilities in one tensor
      label_batch
```

```
[14]: <tf.Tensor: shape=(32,), dtype=int32, numpy=
      array([1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0,
             0, 0, 1, 0, 1, 1, 1, 1, 0, 0], dtype=int32)>
```

Now decode the predictions made by the model. Earlier, when you printed the shape of the batch, it would have returned (32, 1000). The number 32 refers to the batch size and 1000 refers to the 1000 classes the model was pretrained on. The predictions returned by the base model below follow this format:

First the class number, then a human-readable label, and last the probability of the image belonging to that class. You'll notice that there are two of these returned for each image in the batch - these are the top two probabilities returned for that image.

```
[15]: base_model.trainable = False
      image_var = tf.Variable(preprocess_input(image_batch))
      pred = base_model(image_var)

      # Function to decode predictions
      def decode_predictions(preds, top=2):
          results = []
          for pred in preds:
              top_indices = pred.argsort()[-top:][::-1]
              result = [tuple(class_index[str(i)]) + (pred[i],) for i in top_indices]
              results.append(result)
          return results

      with open("imagenet_base_model/imagenet_class_index.json", 'r') as f:
          class_index = json.load(f)
```

```

decoded_predictions = decode_predictions(pred.numpy(), top=2)
print(decoded_predictions)

```

```

[(['n02489166', 'proboscis_monkey', 0.10329965), ('n02102177',
'Welsh_springer_spaniel', 0.07883611)], [(['n02125311', 'cougar', 0.1654676),
('n02389026', 'sorrel', 0.10764261)], [(['n02437312', 'Arabian_camel',
0.2923283), ('n02437616', 'llama', 0.27713484)], [(['n03944341', 'pinwheel',
0.31154886), ('n03047690', 'clog', 0.052500293)], [(['n02454379', 'armadillo',
0.73107153), ('n01990800', 'isopod', 0.038719974)], [(['n02437312',
'Arabian_camel', 0.25663644), ('n02422106', 'hartebeest', 0.12122728)],
[(['n02437616', 'llama', 0.6612557), ('n02090721', 'Irish_wolfhound',
0.23782855)], [(['n02133161', 'American_black_bear', 0.82735676), ('n02134418',
'sloth_bear', 0.02925945)], [(['n01518878', 'ostrich', 0.9267562), ('n02002724',
'black_stork', 0.0017766367)], [(['n01518878', 'ostrich', 0.94954586),
('n02018795', 'bustard', 0.0028661634)], [(['n02437616', 'llama', 0.8699833),
('n02412080', 'ram', 0.076757126)], [(['n02415577', 'bighorn', 0.2429446),
('n02412080', 'ram', 0.160565)], [(['n02437616', 'llama', 0.9473245),
('n02480495', 'orangutan', 0.0076571796)], [(['n09428293', 'seashore',
0.48092392), ('n09421951', 'sandbar', 0.26179993)], [(['n02437312',
'Arabian_camel', 0.95963204), ('n02504458', 'African_elephant', 0.0009881927)],
[(['n02509815', 'lesser_panda', 0.9096807), ('n02443114', 'polecat',
0.014759211)], [(['n01518878', 'ostrich', 0.74165), ('n02002724', 'black_stork',
0.07205889)], [(['n02437312', 'Arabian_camel', 0.49920738), ('n02412080', 'ram',
0.11842591)], [(['n01518878', 'ostrich', 0.87967354), ('n02018795', 'bustard',
0.0077298395)], [(['n02437616', 'llama', 0.82569915), ('n02437312',
'Arabian_camel', 0.010480011)], [(['n01518878', 'ostrich', 0.9612779),
('n02410509', 'bison', 0.0013086519)], [(['n02437616', 'llama', 0.636178),
('n02412080', 'ram', 0.058401026)], [(['n02437616', 'llama', 0.5928003),
('n02417914', 'ibex', 0.039721698)], [(['n02437616', 'llama', 0.83541703),
('n02104029', 'kuvasz', 0.048998024)], [(['n03042490', 'cliff_dwelling',
0.3091509), ('n04208210', 'shovel', 0.06726616)], [(['n02093647',
'Bedlington_terrier', 0.4338772), ('n02113799', 'standard_poodle', 0.4069308)],
[(['n02133161', 'American_black_bear', 0.97880507), ('n02132136', 'brown_bear',
0.0055297976)], [(['n01518878', 'ostrich', 0.83605814), ('n02018795', 'bustard',
0.004823002)], [(['n02133161', 'American_black_bear', 0.9362426), ('n02134418',
'sloth_bear', 0.007733786)], [(['n03240683', 'drilling_platform', 0.04555222),
('n04146614', 'school_bus', 0.033719867)], [(['n02437616', 'llama', 0.9278842),
('n02098286', 'West_Highland_white_terrier', 0.0057286685)], [(['n02437616',
'llama', 0.94477594), ('n02423022', 'gazelle', 0.0054335156)]]

```

Uh-oh. There’s a whole lot of labels here, some of them hilariously wrong, but none of them say “alpaca.”

This is because MobileNet pretrained over ImageNet doesn’t have the correct labels for alpacas, so when you use the full model, all you get is a bunch of incorrectly classified images.

Fortunately, you can delete the top layer, which contains all the classification labels, and create a new classification layer.

3.2 - Layer Freezing with the Functional API

In the next sections, you'll see how you can use a pretrained model to modify the classifier task so that it's able to recognize alpacas. You can achieve this in three steps:

1. Delete the top layer (the classification layer)
 - Set `include_top` in `base_model` as `False`
2. Add a new classifier layer
 - Train only one layer by freezing the rest of the network
 - As mentioned before, a single neuron is enough to solve a binary classification problem.
3. Freeze the base model and train the newly-created classifier layer
 - Set `base_model.trainable=False` to avoid changing the weights and train *only* the new layer
 - Set training in `base_model` to `False` to avoid keeping track of statistics in the batch norm layer

Exercise 2 - alpaca_model

```
[16]: # UNQ_C2
# GRADED FUNCTION
def alpaca_model(image_shape=IMG_SIZE, data_augmentation=data_augmenter()):
    ''' Define a tf.keras model for binary classification out of the
    ↳ MobileNetV2 model
    Arguments:
        image_shape -- Image width and height
        data_augmentation -- data augmentation function
    Returns:
        Returns:
            tf.keras.model
    '''

    input_shape = image_shape + (3,)

    ### START CODE HERE

    base_model_path="imagenet_base_model/
    ↳ without_top_mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_160_no_top.
    ↳ h5"

    base_model = tf.keras.applications.MobileNetV2(input_shape=input_shape,
                                                    include_top=False, # ↳
    ↳ Important!!!!
                                                    weights=base_model_path)

    # freeze the base model by making it non trainable
    base_model.trainable = False

    # create the input layer (Same as the imageNetv2 input size)
```

```

inputs = tf.keras.Input(shape=input_shape)

# apply data augmentation to the inputs
x = data_augmentation(inputs)

# data preprocessing using the same weights the model was trained on
x = preprocess_input(x)

# set training to False to avoid keeping track of statistics in the batch
→norm layer
x = base_model(x, training=False)

# add the new Binary classification layers
# use global avg pooling to summarize the info in each channel
x = tf.keras.layers.GlobalAveragePooling2D()(x)
# include dropout with probability of 0.2 to avoid overfitting
x = tf.keras.layers.Dropout(0.2)(x)

# use a prediction layer with one neuron (as a binary classifier only needs
→one)
outputs = tf.keras.layers.Dense(1)(x)

### END CODE HERE

model = tf.keras.Model(inputs, outputs)

return model

```

Create your new model using the `data_augmentation` function defined earlier.

```
[17]: model2 = alpaca_model(IMG_SIZE, data_augmentation)
```

```

[18]: from test_utils import summary, comparator

alpaca_summary = [['InputLayer', [(None, 160, 160, 3)], 0],
                  ['Sequential', (None, 160, 160, 3), 0],
                  ['TensorFlowOpLayer', [(None, 160, 160, 3)], 0],
                  ['TensorFlowOpLayer', [(None, 160, 160, 3)], 0],
                  ['Functional', (None, 5, 5, 1280), 2257984],
                  ['GlobalAveragePooling2D', (None, 1280), 0],
                  ['Dropout', (None, 1280), 0, 0.2],
                  ['Dense', (None, 1), 1281, 'linear']] #linear is the
→default activation

comparator(summary(model2), alpaca_summary)

for layer in summary(model2):

```

```
print(layer)
```

All tests passed!

```
['InputLayer', [(None, 160, 160, 3)], 0]
['Sequential', (None, 160, 160, 3), 0]
['TensorFlowOpLayer', [(None, 160, 160, 3)], 0]
['TensorFlowOpLayer', [(None, 160, 160, 3)], 0]
['Functional', (None, 5, 5, 1280), 2257984]
['GlobalAveragePooling2D', (None, 1280), 0]
['Dropout', (None, 1280), 0, 0.2]
['Dense', (None, 1), 1281, 'linear']
```

The base learning rate has been set for you, so you can go ahead and compile the new model and run it for 5 epochs:

```
[19]: base_learning_rate = 0.001
      model2.compile(optimizer=tf.keras.optimizers.Adam(lr=base_learning_rate),
                    loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),
                    metrics=['accuracy'])
```

```
[20]: initial_epochs = 5
      history = model2.fit(train_dataset, validation_data=validation_dataset,
                          epochs=initial_epochs)
```

Epoch 1/5

9/9 [=====] - 8s 936ms/step - loss: 0.7851 - accuracy: 0.5496 - val_loss: 0.5759 - val_accuracy: 0.5538

Epoch 2/5

9/9 [=====] - 7s 824ms/step - loss: 0.6459 - accuracy: 0.6298 - val_loss: 0.5222 - val_accuracy: 0.5846

Epoch 3/5

9/9 [=====] - 7s 801ms/step - loss: 0.5379 - accuracy: 0.6450 - val_loss: 0.3952 - val_accuracy: 0.6769

Epoch 4/5

9/9 [=====] - 7s 814ms/step - loss: 0.4970 - accuracy: 0.7328 - val_loss: 0.4095 - val_accuracy: 0.6615

Epoch 5/5

9/9 [=====] - 8s 836ms/step - loss: 0.4533 - accuracy: 0.7366 - val_loss: 0.3209 - val_accuracy: 0.7692

Plot the training and validation accuracy:

```
[21]: acc = [0.] + history.history['accuracy']
      val_acc = [0.] + history.history['val_accuracy']

      loss = history.history['loss']
      val_loss = history.history['val_loss']
```




```

plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),1])
plt.title('Training and Validation Accuracy')

plt.subplot(2, 1, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.ylabel('Cross Entropy')
plt.ylim([0,1.0])
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()

```



```
[22]: class_names
```

```
[22]: ['alpaca', 'not alpaca']
```

The results are ok, but could be better. Next, try some fine-tuning.

3.3 - Fine-tuning the Model

You could try fine-tuning the model by re-running the optimizer in the last layers to improve accuracy. When you use a smaller learning rate, you take smaller steps to adapt it a little more closely to the new data. In transfer learning, the way you achieve this is by unfreezing the layers at the end of the network, and then re-training your model on the final layers with a very low learning rate. Adapting your learning rate to go over these layers in smaller steps can yield more fine details - and higher accuracy.

The intuition for what's happening: when the network is in its earlier stages, it trains on low-level features, like edges. In the later layers, more complex, high-level features like wispy hair or pointy ears begin to emerge. For transfer learning, the low-level features can be kept the same, as they have common features for most images. When you add new data, you generally want the high-level features to adapt to it, which is rather like letting the network learn to detect features more related to your data, such as soft fur or big teeth.

To achieve this, just unfreeze the final layers and re-run the optimizer with a smaller learning rate, while keeping all the other layers frozen.

Where the final layers actually begin is a bit arbitrary, so feel free to play around with this number a bit. The important takeaway is that the later layers are the part of your network that contain the fine details (pointy ears, hairy tails) that are more specific to your problem.

First, unfreeze the base model by setting `base_model.trainable=True`, set a layer to fine-tune from, then re-freeze all the layers before it. Run it again for another few epochs, and see if your accuracy improved!

Exercise 3

```
[23]: # UNQ_C3
base_model = model2.layers[4]
base_model.trainable = True
# Let's take a look to see how many layers are in the base model
print("Number of layers in the base model: ", len(base_model.layers))

# Fine-tune from this layer onwards
fine_tune_at = 120

### START CODE HERE

# Freeze all the layers before the `fine_tune_at` layer
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable = False

# Define a BinaryCrossentropy loss function. Use from_logits=True
loss_function=tf.keras.losses.BinaryCrossentropy(from_logits=True)

# Define an Adam optimizer with a learning rate of 0.1 * base_learning_rate
optimizer = tf.keras.optimizers.Adam(learning_rate=0.1 * base_learning_rate)

# Use accuracy as evaluation metric
metrics=['accuracy']

### END CODE HERE

model2.compile(loss=loss_function,
```

```
optimizer = optimizer,  
metrics=metrics)
```

Number of layers in the base model: 155

```
[24]: assert type(loss_function) == tf.python.keras.losses.BinaryCrossentropy, "Not_  
      ↳the correct layer"  
assert loss_function.from_logits, "Use from_logits=True"  
assert type(optimizer) == tf.keras.optimizers.Adam, "This is not an Adam_  
      ↳optimizer"  
assert optimizer.lr == base_learning_rate / 10, "Wrong learning rate"  
assert metrics[0] == 'accuracy', "Wrong metric"  
  
print('\033[92mAll tests passed!')
```

All tests passed!

```
[25]: fine_tune_epochs = 5  
total_epochs = initial_epochs + fine_tune_epochs  
  
history_fine = model2.fit(train_dataset,  
                           epochs=total_epochs,  
                           initial_epoch=history.epoch[-1],  
                           validation_data=validation_dataset)
```

Epoch 5/10

9/9 [=====] - 10s 1s/step - loss: 0.5864 - accuracy:
0.7252 - val_loss: 0.2713 - val_accuracy: 0.9538

Epoch 6/10

9/9 [=====] - 9s 990ms/step - loss: 0.3833 - accuracy:
0.7939 - val_loss: 0.3108 - val_accuracy: 0.9231

Epoch 7/10

9/9 [=====] - 9s 989ms/step - loss: 0.2987 - accuracy:
0.8664 - val_loss: 0.1653 - val_accuracy: 0.9385

Epoch 8/10

9/9 [=====] - 9s 988ms/step - loss: 0.2381 - accuracy:
0.8969 - val_loss: 0.1014 - val_accuracy: 0.9692

Epoch 9/10

9/9 [=====] - 9s 957ms/step - loss: 0.1664 - accuracy:
0.9313 - val_loss: 0.0989 - val_accuracy: 0.9692

Epoch 10/10

9/9 [=====] - 9s 990ms/step - loss: 0.1561 - accuracy:
0.9427 - val_loss: 0.0777 - val_accuracy: 0.9692


Ahhh, quite an improvement! A little fine-tuning can really go a long way.

```
[26]: acc += history_fine.history['accuracy']  
val_acc += history_fine.history['val_accuracy']
```

```
loss += history_fine.history['loss']
val_loss += history_fine.history['val_loss']
```

```
[27]: plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.ylim([0, 1])
plt.plot([initial_epochs-1, initial_epochs-1],
         plt.ylim(), label='Start Fine Tuning')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(2, 1, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.ylim([0, 1.0])
plt.plot([initial_epochs-1, initial_epochs-1],
         plt.ylim(), label='Start Fine Tuning')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()
```



What you should remember:

- To adapt the classifier to new data: Delete the top layer, add a new classification layer, and train only on that layer
- When freezing layers, avoid keeping track of statistics (like in the batch normalization layer)
- Fine-tune the final layers of your model to capture high-level details near the end of the network and potentially improve accuracy

1.3 Congratulations!

You've completed this assignment on transfer learning and fine-tuning. Here's a quick recap of all you just accomplished:

- Created a dataset from a directory
- Augmented data with the Sequential API
- Adapted a pretrained model to new data with the Functional API and MobileNetV2
- Fine-tuned the classifier's final layers and boosted the model's accuracy

That's awesome!